

# CS 343 - Operating Systems

## Module-2E

### Introduction to Threads



**Dr. John Jose**

**Associate Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati**

# Session Outline

- ❖ **Process vs Threads**
- ❖ **Thread model**
- ❖ **Multithreaded programs**
- ❖ **User and Kernel threads**
- ❖ **Multithread mapping models**
- ❖ **Thread libraries and operations**

# Concept of Threads

- ❖ Thread is a flow of control within a process.
  - ❖ single-threaded process, multi-threaded process.
- ❖ It is a basic unit of CPU utilization, which comprise
  - ❖ a thread ID, program counter, register set, stack.
- ❖ Shares with other threads belonging to the same process its code section, data section, and other OS resources (open files and signal)
- ❖ If a process has multiple threads of control, it can perform more than one task at a time.

# The Thread Model


- ❖ Items shared by all threads in a process
- ❖ Items private to each thread

## Per process items



Address space  
Global variables  
Open files  
Child processes  
Pending alarms  
Signals and signal handlers  
Accounting information

## Per thread items

Program counter  
Registers  
Stack  
State



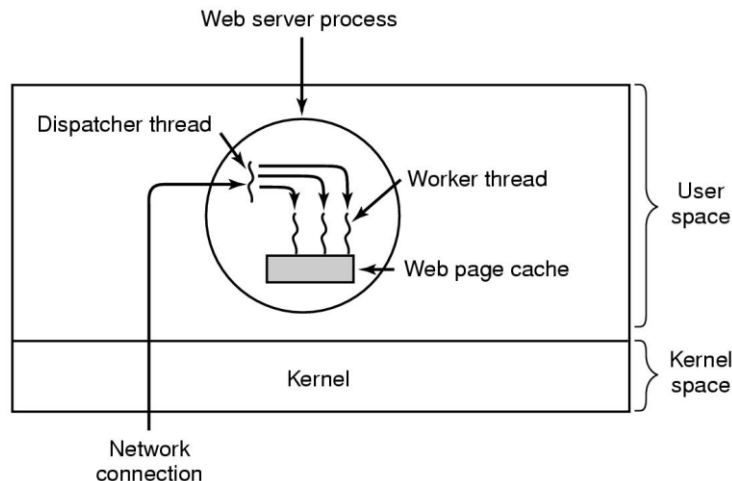
# The Thread Model



Three processes each with one thread   Vs   One process with three threads


# Multi-threaded programs

- ❖ Many software packages that run on modern OS are multi-threaded.
- ❖ A web browser might have
  - ❖ One thread display images or text
  - ❖ Another thread retrieves data from the network



# Multi-threaded programs

- ❖ Many software packages that run on modern OS are multi-threaded.
- ❖ A word processor may have
  - ❖ A thread for displaying graphics
  - ❖ Another thread for responding to keystrokes from the user
  - ❖ A third thread for performing spelling and grammar checking



# Multi-threaded programs

## ❖ Types of Web Server

- ❖ **Single-threaded web server**: a client might have to wait for its request to be serviced.
- ❖ **Multi-threaded web server**: less overhead in thread creation, concurrent service to multiple client.

## ❖ Many OS kernels are now multi-threaded

- ❖ Several threads operates in the kernel
- ❖ Each thread performs a specific task, such as managing devices or interrupt handling.



# Benefits of multi-threaded programming

## ❖ Responsiveness

- ❖ Multithreading an interactive application may allow a program to continue running even if part of it is blocked or doing a lengthy operation.

## ❖ Resource Sharing

- ❖ Threads share the memory and the resources of the process to which they belong.

## ❖ Economy

- ❖ Because threads in a process share the resources, it is more economical to create and context-switch threads.

## ❖ Utilization of Multi-Processor Architectures

- ❖ Threads may be running in parallel on different processors.

# Two types of threads

## ❖ User Thread

- ❖ User-level thread are threads that are visible to the programmer and are unknown to the kernel.
- ❖ User thread are supported above the kernel and are managed without kernel support.
- ❖ Thread management done by [user-level threads library](#)
- ❖ [Three](#) primary thread libraries:
  - ❖ [POSIX](#) Pthreads
  - ❖ [Win32](#) threads
  - ❖ [Java](#) threads

# Two types of threads


## ❖ **Kernel Thread**

- ❖ OS kernel supports and manages kernel-level threads
- ❖ The threads are supported and managed directly by the operating system.


## ❖ Examples

- ❖ Windows 10
- ❖ Solaris
- ❖ Linux
- ❖ Tru64 UNIX
- ❖ Mac OS X

# Implementing Threads in User Space



A user-level threads package




A threads package managed by the kernel

# Multithreading Models

- ❖ A Relationship between **user threads** and **kernel threads**.
  - ❖ Many-to-One
  - ❖ One-to-One
  - ❖ Many-to-Many
  - ❖ Two Level Model

# Many-to-One

- ❖ Many user-level threads mapped to single kernel thread
  - ❖ Thread management is done by the thread library in user space
  - ❖ Can create as many user threads as you wish.
  - ❖ The entire process will block when a thread makes a blocking system call.
  - ❖ Even on multiprocessors, threads are unable to run in parallel
- ❖ Examples:
  - ❖ Solaris Green Threads
  - ❖ GNU Portable Threads




# One-to-One

- ❖ Each user-level thread maps to a kernel thread
  - ❖ Provides more concurrency than the many-to-one model
  - ❖ Allows another thread to run when a thread is in blocking system call
  - ❖ Creating a user thread requires creating the corresponding kernel thread. (overhead)
  - ❖ The number of threads a process can create is smaller than many-to-one model. (careful not to create too many thread)


- ❖ Examples

- ❖ Windows NT/XP/2000
- ❖ Linux
- ❖ Solaris 9 and later



# Many-to-Many Model


- ❖ Allows many user level threads to be mapped to smaller or equal kernel threads
  - ❖ Allows the OS to create a sufficient number of kernel threads
  - ❖ The number of kernel threads may be specific to either a application or machine
- ❖ Examples
  - ❖ Solaris prior to version 9
  - ❖ Windows NT/2000 with the ThreadFiber package





# Two-Level Model

- ❖ One popular variation on many-to-many model
  - ❖ Similar to Many-to-Many model,
  - ❖ Many user-level threads are multiplexed to a smaller or equal number of kernel threads
  - ❖ But it allows a user thread to be **bound** to a kernel thread
- ❖ Examples
  - ❖ IRIX
  - ❖ HP-UX
  - ❖ Tru64 UNIX
  - ❖ Solaris 8 and earlier



# Thread Libraries



- ❖ A **thread library** provides the programmer an API for creating and managing threads
- ❖ Two primary ways to implement
  - ❖ **to provide a library entirely in user space with no kernel support.**
    - ❖ All code and data structures for the library exist in user space
    - ❖ Every function call executes in user mode, not in kernel mode
  - ❖ **to implement a kernel-level library supported by OS**
    - ❖ All code and data structures exist in kernel space
    - ❖ Invoking functions result in a system call to the kernel

# Thread Libraries

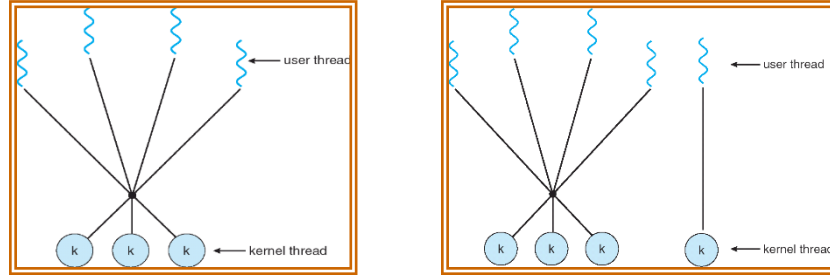
- ❖ A **thread library** provides the programmer an API for creating and managing threads
- ❖ **Three** main thread libraries
  - ❖ **POSIX Pthreads** - Solaris, Linux, Mac OS, Tru64 UNIX
  - ❖ **Win32 Thread** - Windows
  - ❖ **Java Thread** - Java

# Light Weight Processes (LWP)

- ❖ Most popular mapping model - **many-to-many** or **two-level** mode
- ❖ Light Weight Process (LWP)
  - ❖ An intermediate data structure between user and kernel threads
  - ❖ A user thread is attached to a LWP
  - ❖ Each LWP is attached to a kernel thread
  - ❖ OS schedules the kernel threads (not processes) to run on CPU
  - ❖ If a kernel thread blocks, LWP blocks, and the user thread blocks



# Light Weight Processes (LWP)



- ❖ To the user thread library, LWP appears to be a virtual CPU on which the application can schedule a user thread to run.
- ❖ The user thread library is responsible for scheduling among user threads to the LWPs. It is similar to the CPU scheduling in kernel.
- ❖ In general, context switching between user threads involves taking a user thread of its LWP and replacing it with another thread.

# Threading Issues

- ❖ Semantics of `fork()` and `exec()` system calls
- ❖ Thread cancellation
- ❖ Signal handling
- ❖ Thread pools
- ❖ Scheduler activations

# Semantics of fork() and exec()

- ❖ **fork()** - system call is used to create a duplicate process.
- ❖ **exec()** - system call is used to create a new separate process.
- ❖ Fork() starts a new process which is a copy of the one that calls it, while exec () replaces the current process image with another new one.
- ❖ Both parent and child processes are executed simultaneously in case of fork()
- ❖ In exec() control never returns to the original program unless there is an exec() error.

# Thread Cancellation

- ❖ Terminating a thread before it has finished by other threads
  - ❖ Ex, multiple threads search DB, one thread returns result. The remaining thread might be canceled.
- ❖ Two general approaches to cancel the target thread
  - ❖ **Asynchronous cancellation** terminates the target thread immediately
  - ❖ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled



# Thread Cancellation

## ❖ Asynchronous cancellation

- ❖ The difficulty with cancellation occurs in situation where
  - ❖ resources have been allocated to a canceled thread, or
  - ❖ where a thread is canceled while in the midst of updating data it is sharing with other threads

## ❖ Deferred cancellation

- ❖ One thread indicates that target thread is to be canceled.
- ❖ Cancellation occurs only after the target thread has checked a flag to determine if it should be canceled or not

# Signal Handling

- ❖ **Signals** are used in UNIX systems to notify a process that a particular event has occurred
- ❖ **Two types of signals**
  - ❖ Synchronous signals [illegal memory access, division by 0]
  - ❖ Asynchronous signals [Specific keystrokes (Ctrl-C), timer expire]
- ❖ What happen when a signal generated?
  - ❖ Signal is generated by particular event
  - ❖ Signal is delivered from kernel to a process
  - ❖ Signal is handled

# Signal Handling

- ❖ Every signal may be handled by one of two possible handlers.
  - ❖ A default signal handler
  - ❖ A user-defined signal handler
- ❖ Every signal has a default signal handler that is run by the kernel
- ❖ The default action can be overridden by a user-defined signal handler
- ❖ Options when a signal occurs on a multi-threaded process
  - ❖ Deliver the signal to the thread to which the signal applies
  - ❖ Deliver the signal to every thread in the process
  - ❖ Deliver the signal to certain threads in the process
  - ❖ Assign a specific thread to receive all signals for the process

# Thread Pools

- ❖ **Pool** of threads where they await work
- ❖ A process creates few threads at start up and place into a pool
- ❖ When receiving a request, server awakens a thread from the pool and passes the request to service
- ❖ Once the thread completes its service, it returns to the pool and await more work.
- ❖ If the pool contains no available thread, the server waits until one becomes free.
- ❖ Thread pools are faster to service a request with an existing thread than create a new thread
- ❖ Allows the number of threads in the application(s) to be bound to the size of the pool

# Scheduler Activations

- ❖ Both Many-to-Many and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ❖ Scheduler activations provide upcalls - a communication mechanism from the kernel to the thread library
- ❖ This communication allows an application to maintain the correct number kernel threads

# Summary

- ❖ A thread is a flow of control within process.
- ❖ Multithreaded process contains several different flows of control within the same address space.
- ❖ Benefits of multi-threading includes
  - ❖ Increased responsiveness, Resource sharing within the process
  - ❖ Economy, Ability to take advantage of multiprocessor architecture
- ❖ User-level thread are thread that are visible to the programmer and are unknown to the kernel.

# Summary

- ❖ OS kernel supports and manages kernel-level threads
- ❖ Three types of models relates user and kernel threads
  - ❖ One-to-one, many-to-one, many-to-many
- ❖ Thread libraries provide the application programmer with an API for creating and managing threads
  - ❖ POSIX Pthreads, Win32 threads, Java threads
- ❖ Multithreaded programs introduces several issues
  - ❖ fork()/exec(), thread cancellation, signal handling, thread pools and schedule activation.



**johnjose@iitg.ac.in**

**<http://www.iitg.ac.in/johnjose/>**