

# CS 343 - Operating Systems

## Module-4D

### Virtual Memory Techniques



**Dr. John Jose**

**Associate Professor**

**Department of Computer Science & Engineering**

**Indian Institute of Technology Guwahati**

# Overview of Memory Management

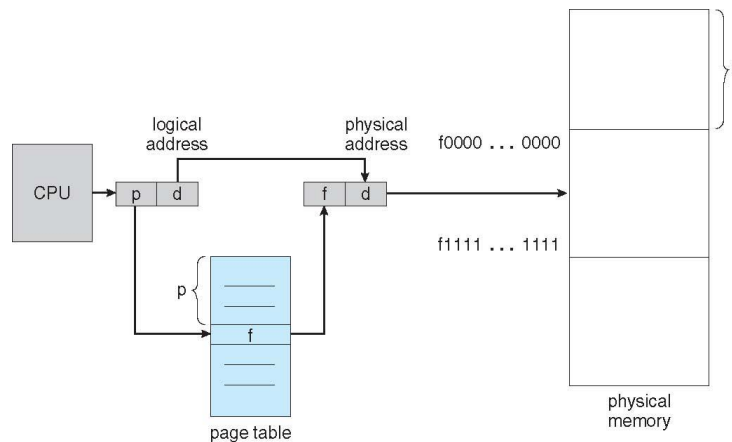
- ❖ Demand Paging
- ❖ Copy-on-Write
- ❖ Page Replacement
- ❖ Allocation of Frames
- ❖ Thrashing
- ❖ Memory-Mapped Files

# Background


- ❖ Program needs to be in memory to execute, but entire program rarely used
- ❖ Entire program code not needed at same time
- ❖ Consider ability to execute partially-loaded program
  - ❖ Program no longer constrained by limits of physical memory
  - ❖ Each program takes less memory while running → more programs run at the same time
  - ❖ Increased CPU utilization and throughput with no increase in response time or turnaround time
  - ❖ Logical address space can therefore be much larger than physical address space

# Background

- ❖ **Virtual address space** – logical view of how process is stored in memory
  - ❖ Usually start at address 0, contiguous addresses until end of space
  - ❖ Meanwhile, physical memory organized in page frames
  - ❖ MMU must map logical to physical
- ❖ Virtual memory can be implemented via:
  - ❖ Demand paging
  - ❖ Demand segmentation




# Virtual Memory To Physical Memory Mapping





# Demand Paging

- ❖ Bring a page into memory only when it is needed
  - ❖ Less I/O needed, no unnecessary I/O
  - ❖ Less memory needed
  - ❖ Faster response
  - ❖ More users
- ❖ **Lazy swapper** – never swaps a page into memory unless page will be needed



# Valid-Invalid Bit

- ❖ With each page table entry a valid–invalid bit is associated (**v**  $\Rightarrow$  in-memory – **memory resident**, **i**  $\Rightarrow$  not-in-memory)
- ❖ Initially valid–invalid bit is set to **i** on all entries




Page Table When Some Pages Are Not in Main Memory

# Page Fault

- ❖ If there is a reference to a page and if it is not mapped to a frame, it will generate a trap interrupt to operating system: **page fault**
- ❖ **Page fault:** page not found in main memory
- ❖ Find free frame (victim frame)
- ❖ Swap page into frame via scheduled disk operation
- ❖ Reset tables to indicate page now in memory: Set validation bit = **1**
- ❖ Restart the instruction that caused the page fault



# Steps in Handling a Page Fault



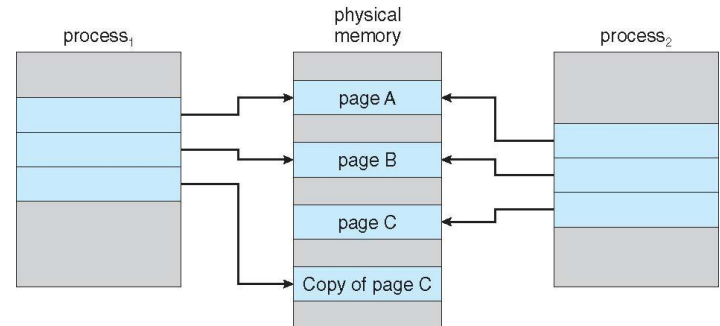

# Demand Paging Overhead

- ❖ Page Fault Rate  $0 \leq p \leq 1$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault
- ❖ Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access} + p (\text{page fault overhead} \\ + \text{swap page out} + \text{swap page in})$$

# Copy-on-Write


- ❖ **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
- ❖ If either process modifies a shared page, only then is the page copied
- ❖ COW allows more efficient process creation as only modified pages are copied
- ❖ In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages



# Basic Page Replacement

- ❖ Find the location of the desired page on disk
- ❖ Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**
  - Write victim frame to disk if dirty
- ❖ Bring the desired page into the (newly) free frame; update the page and frame tables
- ❖ Continue the process by restarting the instruction that caused the trap

# Page Replacement




# Page Replacement Algorithm

## ❖ Page-replacement algorithm

- ❖ Want lowest page-fault rate on both first access and re-access
- ❖ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
  - ❖ String is just page numbers, not full addresses
  - ❖ Repeated access to the same page does not cause a page fault
  - ❖ **FIFO, LIFO,**
  - ❖ **Optimal,**
  - ❖ **LRU, LRU approximations,**
  - ❖ **LFU, MFU**

# Page Faults Vs Number of Frames



# First-In-First-Out (FIFO) Algorithm

- ❖ Ref. string: **7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**
- ❖ 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2	2	4	4	4	0		0	0		7	7	7
	0	0	0		3	3	3	2	2	2		1	1		1	0	0
		1	1		1	0	0	0	3	3		3	2		2	2	1

page frames


15 page faults

- ❖ How to track ages of pages? - Use a FIFO queue



# Belady's Anomaly

- ❖ Consider 1,2,3,4,1,2,5,1,2,3,4,5
- ❖ Adding more frames can cause more page faults! - **Belady's Anomaly**



# Optimal Algorithm

- ❖ Replace page that will not be used for longest period of time
- ❖ Practical difficulty- Can't read the future
- ❖ Used for measuring how well your algorithm performs

reference string


7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2				2				2			2				7
	0	0	0				0				0			0				0
		1	1				3				3			1				1

page frames

# Least Recently Used (LRU) Algorithm

- ❖ Use past knowledge rather than future
- ❖ Replace page that has not been used in the most amount of time
- ❖ Associate time of last use with each page



- ❖ 12 faults – better than FIFO but worse than OPT
- ❖ Generally good algorithm and frequently used

# LRU Algorithm Implementation

## ❖ Counter implementation

- ❖ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- ❖ When a page needs to be changed, look at the counters to find smallest value.

# LRU Approximation Algorithms


## ❖ Reference bit

- ❖ With each page associate a bit, initially = 0
- ❖ When page is referenced, bit set to 1
- ❖ Replace any with reference bit = 0 (if one exists)


## ❖ Second-chance algorithm

- ❖ If page to be replaced has
  - ❖ Reference bit = 0 → replace it
  - ❖ reference bit = 1 then:
    - ❖ set reference bit 0, leave page in memory
    - ❖ replace next page, subject to same rules

# Second-Chance Page-Replacement



(a)



(b)

# Enhanced Second-Chance Algorithm

- ❖ Improve algorithm by using reference bit and modify bit
- ❖ Take ordered pair (reference, modify)
- ❖ (0, 0) neither recently used nor modified – best page to replace
- ❖ (0, 1) not recently used but modified – not quite as good, must write out before replacement
- ❖ (1, 0) recently used but clean – probably will be used again soon
- ❖ (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement

# Counting Algorithms

- ❖ Keep a counter of the number of references that have been made to each
- ❖ **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- ❖ **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used



# Page-Buffering Algorithms

- ❖ Always keep a pool of free frames
  - ❖ When needed, frame is always available, not found at fault time
  - ❖ Read page into free frame and select victim to evict and add to free pool
  - ❖ When convenient, evict victim
- ❖ Keep list of modified pages
  - ❖ When free, write to backing store and set to non-dirty
- ❖ Possibly, keep free frame contents intact and note what is in them
  - ❖ If referenced again before reused, no need to load contents again from disk
  - ❖ Generally useful to reduce penalty if wrong victim frame selected



**johnjose@iitg.ac.in**

**<http://www.iitg.ac.in/johnjose/>**