

CS_344 Assignment 0A

Akshat Mittal, 200101011

Aug 12, 2022

Part 1: PC Boot Strap

Exercise 1

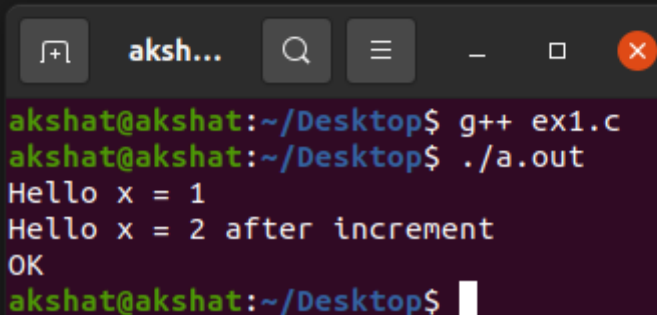
Modified code is shown along with output below:

```
// Simple inline assembly example
#include <stdio.h>
int main(int argc, char **argv){
    int x = 1;
    printf("Hello x = %d\n", x);

    // Put in-line assembly here to increment
    // the value of x by 1 using in-line assembly

    asm("add $1 , %0"
        : "=r"(x)
        : "0"(x)
        );

    printf("Hello x = %d after increment\n",x);
    if(x == 2){
        printf("OK\n");
    }
    else{
        printf("ERROR\n");
    }
}
```



```
akshat@akshat:~/Desktop$ g++ ex1.c
akshat@akshat:~/Desktop$ ./a.out
Hello x = 1
Hello x = 2 after increment
OK
akshat@akshat:~/Desktop$
```

Using **extended inline assembly**, I added the value **1** to input register and stored the result in output register.

The syntax is:

```
asm asm-qualifiers (
    AssemblerTemplate
    : OutputOperands
    : InputOperands
    : Clobbers
    : GotoLabels
);
```

(Here, we do not require Clobbers and GotoLabels.)

The **asm** keyword is a GNU extension.

From the output, we can see that value of **x** has been incremented by **1**.

- After this task, we are required to launch **QEMU** after cloning it.
- We will open two terminals, in first one, run ***make qemu-nox-gdb*** to start QEMU and in second one, from the same directory, run ***gdb*** and type ***source .gdbinit*** at the prompt.

Exercise 2

I executed the ***si*** command some number of times and got following results:

```
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb)
[f000:e062] 0xfe062: jne 0xd241d0b2
0x0000e062 in ?? ()
(gdb)
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb)
[f000:e068] 0xfe068: mov %edx,%ss
0x0000e068 in ?? ()
(gdb)
[f000:e06a] 0xfe06a: mov $0x7000,%sp
0x0000e06a in ?? ()
(gdb)
[f000:e070] 0xfe070: mov $0x7c4,%dx
0x0000e070 in ?? ()
(gdb)
[f000:e076] 0xfe076: jmp 0x5576cf26
0x0000e076 in ?? ()
(gdb)
[f000:cf24] 0xfcf24: cli
0x0000cf24 in ?? ()
(gdb)
```

1. Compare the values at 0xffc8 with one at given CS and IP location.
2. Conditionally jump if result of previous instruction is not equal
3. Set %edx to 0 since taking xor with itself results in 0.
4. Set %ss (**stack segment register**) to %edx, i.e., 0.
5. Set %sp (**stack point register**) to 0x700.
6. Set %dx to 0x7c4.
7. Jump to the specified location.
8. Clear the **interrupt** flag.

In short, various **flags** and **registers** were initialized here.

Part 2: The Boot Loader

Exercise 3

- We have to set a breakpoint at location 0x7c00, i.e., the location where boot sector will be loaded, using **b** command: **b *0x7c00**.
- Then using **c** command, we will continue the execution till next breakpoint.
- Then we have to trace the output of **x/Ni** (here N=15) command in source code **bootasm.S** and disassembly in **bootblock.asm**.

```

(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/15i $eip
=> 0x7c00:    cli
0x7c01:    xor    %eax,%eax
0x7c03:    mov    %eax,%ds
0x7c05:    mov    %eax,%es
0x7c07:    mov    %eax,%ss
0x7c09:    in     $0x64,%al
0x7c0b:    test   $0x2,%al
0x7c0d:    jne    0x7c09
0x7c0f:    mov    $0xd1,%al
0x7c11:    out    %al,$0x64
0x7c13:    in     $0x64,%al
0x7c15:    test   $0x2,%al
0x7c17:    jne    0x7c13
0x7c19:    mov    $0xdf,%al
0x7c1b:    out    %al,$0x60

```

Fig. GDB Output

```

10 .code16 # Assemble for 16-bit mode
11 .globl start
12 start:
13 cli # BIOS enabled interrupts;
14 7c00: fa cli
15
16 # Zero data segment registers DS, ES, and SS.
17 xorw %ax,%ax # Set %ax to zero
18 7c01: 31 c0 xor %eax,%eax
19 movw %ax,%ds # -> Data Segment
20 7c03: 8e d8 mov %eax,%ds
21 movw %ax,%es # -> Extra Segment
22 7c05: 8e c0 mov %eax,%es
23 movw %ax,%ss # -> Stack Segment
24 7c07: 8e d0 mov %eax,%ss
25
26 00007c09 <seta20.1>:
27
28 # Physical address line A20 is tied to zero so that the
29 # with 2 MB would run software that assumed 1 MB. Undo
30 3 references
31 seta20.1:
32 inb $0x64,%al # Wait for not busy
33 7c09: e4 64 in $0x64,%al
34 testb $0x2,%al
35 7c0b: a8 02 test $0x2,%al
36 jnz seta20.1
37 7c0d: 75 fa jne 7c09 <seta20.1>
38 movb $0xd1,%al # 0xd1 -> port 0x64
39 7c0f: b0 d1 mov $0xd1,%al
40 outb %al,$0x64
41 7c11: e6 64 out %al,$0x64
42
43 00007c13 <seta20.2>:
44
45 3 references
46 seta20.2:
47 inb $0x64,%al # Wait for not busy
48 7c13: e4 64 in $0x64,%al

```

Fig. bootblock.asm

```

10 .code16 # Ass
11 .globl start
12 start:
13 cli Fig. bootasm.S # BI
14
15 # Zero data segment registers DS,
16 xorw %ax,%ax # Set
17 movw %ax,%ds # ->
18 movw %ax,%es # ->
19 movw %ax,%ss # ->
20
21 # Physical address line A20 is t
22 # with 2 MB would run software th
23 seta20.1:
24 inb $0x64,%al
25 testb $0x2,%al
26 jnz seta20.1
27
28 movb $0xd1,%al
29 outb %al,$0x64
30
31 seta20.2:
32 inb $0x64,%al
33 testb $0x2,%al
34 jnz seta20.2
35
36 movb $0xdf,%al
37 outb %al,$0x60
38

```

Fig. bootasm.S

- It is observed that all the instructions are same except few changes in syntax (the suffix **b/w/l** stands for **block/ word/ long word** respectively and is not present in output).
- Here, the underlying commands are same, the only difference being the style of writing.
- The line 2 of gdb output is same as line 17 of bootblock.asm and line 16 of bootasm.S (marked with red arrow) and similarly those marked with blue arrows are same.

- Now, we have to trace various parts of source code and corresponding assembly instructions in **bootmain.c** and **bootblock.asm**.

1. readsect(): Read a single sector at offset from disk.

```
// Read a single sector at offset into dst.
void
readsect(void *dst, uint offset)
{
    // Issue command.
    waitdisk();
    outb(0x1F2, 1); // count = 1
    outb(0x1F3, offset);
    outb(0x1F4, offset >> 8);
    outb(0x1F5, offset >> 16);
    outb(0x1F6, (offset >> 24) | 0xE0);
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Fig. Source code in bootmain.c

```
// Read a single sector at offset into dst.
381 references
void
readsect(void *dst, uint offset)
{
    7c90: f3 0f 1e fb      endbr32
    7c94: 55                push    %ebp
    7c95: 89 e5             mov     %esp,%ebp
    7c97: 57                push    %edi
    7c98: 53                push    %ebx
    7c99: 8b 5d 0c          mov     0xc(%ebp),%ebx
    // Issue command.
    waitdisk();
    7c9c: e8 dd ff ff ff    call    7c7e <waitdisk>
}
```

Fig. Corresponding assembly code in bootblock.asm

2. For loop that reads remaining sector of kernel from disk.

```
33 // Load each program segment (ignores ph flags).
34 ph = (struct proghdr*)((uchar*)elf + elf->phoff);
35 eph = ph + elf->phnum;
36 for(; ph < eph; ph++){
37     pa = (uchar*)ph->paddr;
38     readseg(pa, ph->filesz, ph->off);
39     if(ph->memsz > ph->filesz)
40         stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
41 }
42 // Call the entry point from the ELF header.
43 // Does not return!
44 entry = (void(*) (void))(elf->entry);
45 entry();
46 }
47
48
```

Fig. bootmain.c

For loop runs from line 327 – 348 in bootblock.asm. We can see on line 329 loop condition is checked. If the condition fails, on line 330, conditional jump will move to address **0x7d91** on line 319.

Now the control will be given to the OS.

```
315 for(; ph < eph; ph++){
316     7d8d: 39 f3             cmp     %esi,%ebx
317     7d8f: 72 15             jb      7da6 <bootmain+0x5d>
318     entry();
319     7d91: ff 15 18 00 01 00 call    *0x10018
320 }
321     7d97: 8d 65 f4          lea     -0xc(%ebp),%esp
322     7d9a: 5b                pop     %ebx
323     7d9b: 5e                pop     %esi
324     7d9c: 5f                pop     %edi
325     7d9d: 5d                pop     %ebp
326     7d9e: c3                ret
327 for(; ph < eph; ph++){
328     7d9f: 83 c3 20          add     $0x20,%ebx
329     7da2: 39 de             cmp     %ebx,%esi
330     7da4: 76 eb             jbe     7d91 <bootmain+0x48>
331     pa = (uchar*)ph->paddr;
332     7da6: 8b 7b 0c          mov     0xc(%ebx),%edi
333     readseg(pa, ph->filesz, ph->off);
334     7da9: 83 ec 04          sub     $0x4,%esp
335     7dac: ff 73 04          pushl   0x4(%ebx)
336     7daf: ff 73 10          pushl   0x10(%ebx)
337     7db2: 57                push    %edi
338     7db3: e8 44 ff ff ff    call    7cfc <readseg>
339     if(ph->memsz > ph->filesz)
340     7db8: 8b 4b 14          mov     0x14(%ebx),%ecx
341     7dbb: 8b 43 10          mov     0x10(%ebx),%eax
342     7dbe: 83 c4 10          add     $0x10,%esp
343     7dc1: 39 c1             cmp     %eax,%ecx
344     7dc3: 76 da             jbe     7d9f <bootmain+0x56>
345     stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
346     7dc5: 01 c7             add     %eax,%edi
347     7dc7: 29 c1             sub     %eax,%ecx
348 }
349
```

Fig. bootblock.asm

```
(gdb) b *0x7d91
Breakpoint 1 at 0x7d91
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d91:      call    *0x10018

Thread 1 hit Breakpoint 1, 0x00007d91 in ?? ()
```

Setting breakpoint at **0x7d91**, i.e., the end of for loop and stepping through boot loader.

- Now let us answer the questions asked:


```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdt_desc
movl    %cr0, %eax
orl     $CR0_PE, %eax
movl    %eax, %cr0

//PAGEBREAK!
# Complete the transition to 32-bit protected mode by using a long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp    $(SEG_KCODE<<3), $start32
```

1. The **ljmp** command makes instruction pointer move from 16-bit to 32-bit mode. It is caused by the 4 lines above this command which ask OS to switch. All instructions after this are in 32-bit mode.

```
(gdb) b *0x7d91
Breakpoint 1 at 0x7d91
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d91:      call    *0x10018

Thread 1 hit Breakpoint 1, 0x00007d91 in ?? ()
(gdb) si
=> 0x10000c:    mov     %cr4,%eax
0x0010000c in ?? ()
(gdb)
```

2. To do so, we set breakpoint at **0x791** (as found above) and using **c** and **si** command found the required instructions.

Last instruction of boot loader:

*call *0x10018*

First instruction of kernel:

mov %cr4, %eax

```
// Load each program segment (ignores ph flags).
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```

3. The for loop runs starting from **ph** whose value can be found out using **ELF header** file and number of iterations are given using **elf->phnum**. Thus, end of for loop, i.e., **eph** becomes: **eph = ph + elf->phnum**.

Exercise 4

• The first part says to understand the output of *pointer.c* file line by line. Here is the output upon execution of the program:

```
akshat@akshat:~/Desktop$ gcc pointer.c
akshat@akshat:~/Desktop$ ./a.out
1: a = 0x7ffe0788ec40, b = 0x55c02c4602a0, c = 0x7ffe0788ec67
2: a[0] = 200, a[1] = 101, a[2] = 102, a[3] = 103
3: a[0] = 200, a[1] = 300, a[2] = 301, a[3] = 302
4: a[0] = 200, a[1] = 400, a[2] = 301, a[3] = 302
5: a[0] = 200, a[1] = 128144, a[2] = 256, a[3] = 302
6: a = 0x7ffe0788ec40, b = 0x7ffe0788ec44, c = 0x7ffe0788ec41
```

```
int a[4];
int *b = (int*)malloc(16);
int *c;
int i;

printf("1: a = %p, b = %p, c = %p\n", a, b, c);
```

Line 1: The variables **a** and **c** belong to **stack** memory address while **malloc** is used to dynamically allocate memory in the **heap** space. Hence, the address of **b** is at totally different location.

```

c = a;
for (i = 0; i < 4; i++)
a[i] = 100 + i;
c[0] = 200;
printf("2: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
      a[0], a[1], a[2], a[3]);

```

another values **a[1...3]** remain as after the **for loop**.

Line 2: The pointer **c** is made to point the location same as **a**. Thus, after **a[0]** is changed to **100** inside **for loop**, the **c[0]** changes it to **200** while

```

c[1] = 300;
*(c + 2) = 301;
3[c] = 302;
printf("3: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
      a[0], a[1], a[2], a[3]);

```

are other ways of dereferencing and which change values of **a[2]** and **a[3]** respectively.

Line 3: Now **c[1]** is changed to **300** which changes **a[1]** as well since **c** is still storing address of **a**. The next lines

```

c = c + 1;
*c = 400;
printf("4: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
      a[0], a[1], a[2], a[3]);

```

than **a[0]** and hence value of **a[1]** is changed to **400**.

Line 4: **c = c+1** increases the address inside **c** by **4 bytes** due to which **c** now stores the address of **a[1]** rather

```

c = (int *) ((char *) c + 1);
*c = 500;
printf("5: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
      a[0], a[1], a[2], a[3]);

```

specific value in **a** but an **address between a[1] and a[2]** and updating it affect both.

Line 5: This time, address in **c** is increased by only **1 byte** since **char is 1 byte** and thus **c** does not point to any

Note: Numbers are stored in binary representation from left to right inside memory, i.e., left-most bit is **LSB**. For eg. '40' in 16-bit system will be stored as 00010100 00000000 rather than 00000000 00101000.

a[1] = 400 and a[2] = 301

00001001 10000000 00000000 00000000 10110100 10000000 00000000 00000000
 ^a[1] ^ location c is pointing ^a[2]

After updating *c = 500: 4 bytes from the position of pointer **c** will be affected.

00001001 00101111 10000000 00000000 00000000 10000000 00000000 00000000
 ^a[1] ^ location c is pointing ^a[2]

Hence, new value of **a[1] = 128144** and **a[2] = 256**

```

b = (int *) a + 1;
c = (int *) ((char *) a + 1);
printf("6: a = %p, b = %p, c = %p\n", a, b, c);

```

address inside **a** is same as initial while in **b** is **a+4 byte** and in **c** is **a+1 byte**.

Line 6: It stores the **(address of a) + 4 bytes in b** and **(address of a) + 1 byte in c** (like above one) and thus the

- Now we have to solve **K-splice** problem:

What does this program print?

Let us assume that **x** is stored at address **0x7ffdfbf7f00** (mentioned in problem itself).

```
#include <stdio.h>
int main() {
    int x[5];
    printf("%p\n", x);
    printf("%p\n", x+1);
    printf("%p\n", &x);
    printf("%p\n", &x+1);
    return 0;
}
```

Output:

```
0x7ffdfbf7f00
0x7ffdfbf7f04
0x7ffdfbf7f00
0x7ffdfbf7f14
```

We can draw following conclusions:

- Pointers and arrays are different things.
- If **x** is an integer array, then **x + 1** increments the address by **4 bytes**, i.e., **x[1]**.
- &x** is a pointer to the array and **&x + 1** increments the address by **4*n bytes** (where **n** is size of the array).

Part 3: Loading the Kernel

The commands **objdump -h kernel** and **objdump -h bootblock.o** were executed to inspect various code sections in the kernel file and text sections of the boot loader.

```
akshat@akshat:~/xv6-public$ objdump -h kernel
kernel:      file format elf32-i386

Sections:
Idx Name          Size      VMA       LMA       File off  Algn
 0 .text          000070da  80100000  00100000  00001000  2**4
   CONTENTS, ALLOC, LOAD, READONLY, CODE
 1 .rodata        000009cb  801070e0  001070e0  000080e0  2**5
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 2 .data          00002516  80108000  00108000  00009000  2**12
   CONTENTS, ALLOC, LOAD, DATA
 3 .bss           0000af88  8010a520  0010a520  0000b516  2**5
   ALLOC
 4 .debug_line     00006cb5  00000000  00000000  0000b516  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 5 .debug_info     000121ce  00000000  00000000  000121cb  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 6 .debug_abbrev   00003fd7  00000000  00000000  00024399  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 7 .debug_aranges  000003a8  00000000  00000000  00028370  2**3
   CONTENTS, READONLY, DEBUGGING, OCTETS
 8 .debug_str      00000eab  00000000  00000000  00028718  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 9 .debug_loc      0000681e  00000000  00000000  000295c3  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
10 .debug_ranges   00000d08  00000000  00000000  0002fde1  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
11 .comment        0000002b  00000000  00000000  00030ae9  2**0
   CONTENTS, READONLY
```

Fig. objdump output for kernel

```
akshat@akshat:~/xv6-public$ objdump -h bootblock.o
bootblock.o:  file format elf32-i386

Sections:
Idx Name          Size      VMA       LMA       File off  Algn
 0 .text          000001d3  00007c00  00007c00  00000074  2**2
   CONTENTS, ALLOC, LOAD, CODE
 1 .eh_frame       000000b0  00007dd4  00007dd4  00000248  2**2
   CONTENTS, ALLOC, LOAD, READONLY, DATA
 2 .comment        0000002b  00000000  00000000  000002f8  2**0
   CONTENTS, READONLY
 3 .debug_aranges  00000040  00000000  00000000  00000328  2**3
   CONTENTS, READONLY, DEBUGGING, OCTETS
 4 .debug_info     000005d2  00000000  00000000  00000368  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 5 .debug_abbrev   0000022c  00000000  00000000  0000093a  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 6 .debug_line     0000029a  00000000  00000000  00000b66  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 7 .debug_str      00000220  00000000  00000000  00000e00  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 8 .debug_loc      000002bb  00000000  00000000  00001020  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
 9 .debug_ranges   00000078  00000000  00000000  000012db  2**0
   CONTENTS, READONLY, DEBUGGING, OCTETS
```

Fig. objdump output for bootblock.o

- objdump -h** gives the list of all sections in the executable (ELF) file of the specified file, along with some details. It gives us the following information:
 - Size of the section.
 - VMA, i.e. Link Address, of the section, where code execution starts.
 - LMA, i.e. Load Address, of the section, where section is loaded into memory.

Exercise 5

- On tracing through the first few instructions of boot loader, I concluded that the **ljmp** instruction would be the first one to break on changing the link address, i.e., where we transition to 32-bit system.

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdt_desc
movl    %cr0, %eax
orl     $CR0_PE, %eax
movl    %eax, %cr0

//PAGEBREAK!
# Complete the transition to 32-bit protected mode by using a long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp    $(SEG_KCODE<<3), $start32
```

Fig. bootasm.S

- Now, I changed the **link** address in **Makefile** from 0x7c00 to 0x7e00 and then run the **make clean** and **make qemu-nox-gdb** commands.

```
bootblock: bootasm.S bootmain.c
$(CC) $(CFLAGS) -fno-pic -O -nostdinc -I. -c bootmain.c
$(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c bootasm.S
$(LD) $(LDFLAGS) -N -e start -Ttext 0x7E00 -o bootblock.o bootasm.o bootmain.o
$(OBJDUMP) -S bootblock.o > bootblock.asm
$(OBJCOPY) -S -O binary -j .text bootblock.o bootblock
./sign.pl bootblock
```

- Now, I will compare the output of **gdb terminal** in correct link address with that of incorrect link address. Doing so, following conclusions can be drawn:

- All instructions before **ljmp** are same for both the cases.
- There is no message saying “*The target architecture is assumed to be i386*” with incorrect link address which means it does not switch to 32-bit system.
- Instructions after this are different. Hence, **ljmp is first instruction to break**.

```
(gdb)
[ 0:7c25] => 0x7c25: or     $0x1,%ax
0x00007c25 in ?? ()
(gdb)
[ 0:7c29] => 0x7c29: mov    %eax,%cr0
0x00007c29 in ?? ()
(gdb)
[ 0:7c2c] => 0x7c2c: ljmp   $0xb866,$0x87c31
0x00007c2c in ?? ()
(gdb)
The target architecture is assumed to be i386
=> 0x7c31: mov    $0x10,%ax
0x00007c31 in ?? ()
(gdb)
=> 0x7c35: mov    %eax,%ds
0x00007c35 in ?? ()
(gdb)
=> 0x7c37: mov    %eax,%es
0x00007c37 in ?? ()
```

Fig. Output with **correct** link address

```
(gdb)
[ 0:7c25] => 0x7c25: or     $0x1,%ax
0x00007c25 in ?? ()
(gdb)
[ 0:7c29] => 0x7c29: mov    %eax,%cr0
0x00007c29 in ?? ()
(gdb)
[ 0:7c2c] => 0x7c2c: ljmp   $0xb866,$0x87e31
0x00007c2c in ?? ()
(gdb)
[f000:e05b] 0xfe05b: cmpw   $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb)
[f000:e062] 0xfe062: jne    0xd241d0b2
0x0000e062 in ?? ()
(gdb)
[f000:d0b0] 0xfd0b0: cli
0x0000d0b0 in ?? ()
(gdb)
[f000:d0b1] 0xfd0b1: cld
0x0000d0b1 in ?? ()
```

Fig. Output with **incorrect** link address

- From the command **objdump -f kernel**, we can see the entry point, i.e., the start address of kernel is at **0x0010000c**.

```
akshat@akshat:~/xv6-public$ objdump -f kernel

kernel:      file format elf32-i386
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x0010000c
```

Fig. objdump -f kernel

Exercise 6

Following conclusions can be drawn from this experiment:

1. All the **8 words** at the point where BIOS enters the boot-loader (address **0x7c00**) are **0x00000000**.

This is because the address location **0x00100000**, i.e., **1 MB** in the main memory, i.e., the point where the **boot-loader has to load the kernel**. Since the boot-loader's first instruction is yet to be executed after the address **0x7c00**, the **loading of kernel into the RAM has not started** and hence, all the words are having the value **0**.

2. At the second breakpoint i.e., the point where **boot-loader gives the control to kernel**, i.e., **0x7d91** the 8 words now have a **non-zero** value.

This output is expected since by this time the **boot-loader has loaded the kernel** into the memory location starting at address **1 MB**. So the words there contain meaningful information.

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) b *0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d91: call *0x10018

Thread 1 hit Breakpoint 2, 0x00007d91 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
```

- Answer to the **second question** is that second breakpoint was set at the instruction corresponding to the address location **0x7d91**. This is the address of the **entry function** which is responsible for **giving the control to kernel** and is the last instruction to execute of the boot-loader (as already discussed in Exercise 3).