

CS_344 Assignment 1

Group C21: Satvik Tiwari, 200101091

Pranjal Baranwal, 200101083

Akshat Mittal, 200101011

Task: Kernel Threads and Synchronization

Part 1: Kernel Threads

In this part, we have to implement kernel threads using three system calls:

1. *thread_create()*

This will be used to create a new kernel thread which shares same address space. Returns PID of new thread to parent.

Function definition:

*int thread_create(void(*fcn)(void*), void *arg, void*stack)*

```
180 int
181 fork(void)
182 {
183     int i, pid;
184     struct proc *np;
185     struct proc *curproc = myproc();
186
187     // Allocate process.
188     if((np = allocproc()) == 0){
189         return -1;
190     }
191
192     // Copy process state from proc.
193     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
194         kfree(np->kstack);
195         np->kstack = 0;
196         np->state = UNUSED;
197         return -1;
198     }
199     np->sz = curproc->sz;
200     np->parent = curproc;
201     *np->tf = *curproc->tf;
202
203     // Clear %eax so that fork returns 0 in the child.
204     np->tf->eax = 0;
205
206     for(i = 0; i < NOFILE; i++)
207         if(curproc->ofile[i])
208             np->ofile[i] = filedup(curproc->ofile[i]);
209     np->cwd = idup(curproc->cwd);
210
211     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
212     pid = np->pid;
213     acquire(&ptable.lock);
214     np->state = RUNNABLE;
215     release(&ptable.lock);
216     return pid;
217 }
```

Fig. *fork()*

```
527 int
528 thread_create(void (*fcn)(void *), void * arg, void * stack)
529 {
530     if ((uint) stack == 0) {
531         return -1;
532     }
533     int i, pid;
534     struct proc *np;
535     struct proc *curproc = myproc();
536     // Allocate process.
537     if ((np = allocproc()) == 0)
538         return -1;
539     np->pgdir = curproc->pgdir;
540     np->sz = curproc->sz;
541     np->parent = curproc;
542     *np->tf = *curproc->tf;
543     // Mark this proc as a thread
544     np->is_thread = 1;
545     // Clear %eax so that fork returns 0 in the child.
546     np->tf->eax = 0;
547     // set function
548     np->tf->eip = (int) fcn;
549     np->tf->esp = (int) stack + 4096;
550     np->tf->esp -= 4;
551     *((int*) (np->tf->esp)) = (int) arg;
552     np->tf->esp -= 4;
553     *((int*) (np->tf->esp)) = 0xffffffff;
554
555     for (i = 0; i < NOFILE; i++)
556         if (curproc->ofile[i])
557             np->ofile[i] = filedup(curproc->ofile[i]);
558
559     np->cwd = idup(curproc->cwd);
560     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
561     pid = np->pid;
562     acquire(&ptable.lock);
563     np->state = RUNNABLE;
564     release(&ptable.lock);
565     return pid;
566 }
```

Fig. *thread_create()*

We will be using the *fork()* call with some modifications to achieve this. All changes made are as follows:

- Line 530-532: Checks if the allocated stack is not null.
- Line 539: Copy *pgdir* pointer of parent process to child thread; this ensures that address space of both child thread and parent process is same otherwise in *fork()*, it allocates different address space to child process (line 193-198).
- Line 544: Mark this child as a thread process, which will be used later in *thread_join()*.

- Line 548: ***eip*** tells the thread where to start execution from; thus making this to point the function passed as argument.
- Line 549-553: ***esp*** points to top of stack; thus making this to point at top of allocated stack passed as argument. After this, we add ***arg*** and fake PC to the stack which allows thread to exit in case some error occurs during execution.

2. *thread_join()*

This call waits for a child thread that shares the address space with the calling process. Returns PID of waited-for child or -1 if none.

Function definition:

int thread_join(void)

```

265 int
266 wait(void)
267 {
268     struct proc *p;
269     int havekids, pid;
270     struct proc *curproc = myproc();
271
272     acquire(&ptable.lock);
273     for(;;){
274         // Scan through table looking for exited children.
275         havekids = 0;
276         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
277             if(p->parent != curproc)
278                 continue;
279             havekids = 1;
280             if(p->state == ZOMBIE){
281                 // Found one.
282                 pid = p->pid;
283                 kfree(p->kstack);
284                 p->kstack = 0;
285                 freevm(p->pgdir);
286                 p->pid = 0;
287                 p->parent = 0;
288                 p->name[0] = 0;
289                 p->killed = 0;
290                 p->state = UNUSED;
291                 release(&ptable.lock);
292                 return pid;
293             }
294         }
295         // No point waiting if we don't have any children.
296         if(!havekids || curproc->killed){
297             release(&ptable.lock);
298             return -1;
299         }
300         // Wait for children to exit. (See wakeup call in
301         sleep(curproc, &ptable.lock); //DOC: wait-sleep
302     }
303 }

```

Fig. *wait()*

```

568 int
569 thread_join(void)
570 {
571     struct proc *p;
572     int havekids, pid;
573     struct proc *curproc = myproc();
574
575     acquire(&ptable.lock);
576     for(;;){
577         // Scan through table looking for exited children.
578         havekids = 0;
579         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
580             if (p->is_thread == 0 || p->parent != curproc)
581                 continue;
582             havekids = 1;
583             if (p->state == ZOMBIE) {
584                 // Found one.
585                 pid = p->pid;
586                 //kfree(p->kstack);
587                 p->kstack = 0;
588                 //freevm(p->pgdir);
589                 p->pid = 0;
590                 p->parent = 0;
591                 p->name[0] = 0;
592                 p->killed = 0;
593                 p->state = UNUSED;
594                 release(&ptable.lock);
595                 return pid;
596             }
597         }
598         // No point waiting if we don't have any children.
599         if (!havekids || curproc->killed) {
600             release(&ptable.lock);
601             return -1;
602         }
603         // Wait for children to exit. (See wakeup call in pr
604         sleep(curproc, &ptable.lock); //DOC: wait-sleep
605     }
606 }
607 }

```

Fig. *thread_join()*

We will be using the ***wait()*** call with some modifications to achieve this. All changes made are as follows:

- Line 580: We need to check if current process is a thread before continuing.
- Line 586: We do not want to free the stack of the thread since parent process may still use that memory.
- Line 588: We do not want to free the complete page allocated to thread since it was same as that of parent process and parent may still use it. In case of *wait()*, we allocated new space for the child, hence we need to free it.

3. *thread_exit()*

This call allows a thread to terminate.

Function definition:

int thread_exit(void)

```

222 void
223 exit(void)
224 {
225     struct proc *curproc = myproc();
226     struct proc *p;
227     int fd;
228
229     if(curproc == initproc)
230         panic("init exiting");
231
232     // Close all open files.
233     for(fd = 0; fd < NOFILE; fd++){
234         if(curproc->ofile[fd]){
235             fileclose(curproc->ofile[fd]);
236             curproc->ofile[fd] = 0;
237         }
238     }
239
240     begin_op();
241     iput(curproc->cwd);
242     end_op();
243     curproc->cwd = 0;
244     acquire(&ptable.lock);
245
246     // Parent might be sleeping in wait().
247     wakeup1(curproc->parent);
248     // Pass abandoned children to init.
249     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
250         if(p->parent == curproc){
251             p->parent = initproc;
252             if(p->state == ZOMBIE)
253                 wakeup1(initproc);
254         }
255     }
256
257     // Jump into the scheduler, never to return.
258     curproc->state = ZOMBIE;
259     sched();
260     panic("zombie exit");
261 }

```

Fig. `exit()`

```

609 int
610 thread_exit(void)
611 {
612     struct proc *curproc = myproc();
613     struct proc *p;
614     int fd;
615
616     if (curproc == initproc)
617         panic("init exiting");
618
619     // Close all open files.
620     for (fd = 0; fd < NOFILE; fd++) {
621         if (curproc->ofile[fd]) {
622             fileclose(curproc->ofile[fd]);
623             curproc->ofile[fd] = 0;
624         }
625     }
626
627     begin_op();
628     iput(curproc->cwd);
629     end_op();
630     curproc->cwd = 0;
631     acquire(&ptable.lock);
632     // Parent might be sleeping in wait().
633     wakeup1(curproc->parent);
634     // Pass abandoned children to init.
635     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
636         if (p->parent == curproc) {
637             p->parent = initproc;
638             if (p->state == ZOMBIE)
639                 wakeup1(initproc);
640         }
641     }
642     // Jump into the scheduler, never to return.
643     curproc->state = ZOMBIE;
644     sched();
645     panic("zombie exit");
646 }

```

Fig. `thread_exit()`

We will be using ***exit()*** call to achieve this. No changes required here because of a thread is exactly similar to that of a child process. This just passes the control back to parent process.

The above three functions are declared in file **proc.c**. Now we have to create system calls for each of these functions.

First, we make corresponding functions in **sysproc.c** file.

```

93 // Create a new kernel thread
94 int
95 sys_thread_create(void)
96 {
97     // These functions do not take arguments directly
98     // Hence using these lines, we read arguments.
99     void (*fcn)(void*), *arg, *stack;
100     argptr(0, (void*) &fcn, sizeof(void (*)(void *)));
101     argptr(1, (void*) &arg, sizeof(void *));
102     argptr(2, (void*) &stack, sizeof(void *));
103     return thread_create(fcn, arg, stack);
104 }

```

```

106 // Wait for child thread to complete
107 int
108 sys_thread_join(void)
109 {
110     return thread_join();
111 }
112 // Terminate a kernel thread
113 int
114 sys_thread_exit(void)
115 {
116     return thread_exit();
117 }

```

Now, we declare them in various files (similar to all other system calls):

1. `syscall.c`
2. `syscall.h`
3. `user.h`
4. `usys.S`
5. `defs.h`

Now to test our implementation, we will create a file **thread.c** which will make use of our newly created system calls to perform a task. Next, we have to update **Makefile** and include **thread** in UPROGS list to add this as a command line instruction in XV6 kernel which will execute **thread.c**.

```

$ thread
SSttaratrintg idon_gw odrok_:w so:rbk2
: s:b1
Done s:2F9C
Done s:2F78
Threads finished: (13):14, (14):13, shared balance:3207
$ thread
SSttaarrttiinnngg ddoo__wwoorrk:: ss::bb2
1
Done s:2F9C
Done s:2F78
Threads finished: (16):17, (17):16, shared balance:3204
$ thread
SSttaarrtitiinnngg ddoo_w_owrokr: k:s: s:b2b1
Done s:2F9C
Done s:2F78
Threads finished: (19):20, (20):19, shared balance:3203

```

Upon execution of **thread**, we got this output. We can clearly see that **value of total_balance is not equal to expected 6000**, i.e., sum of all account's balance. This happens because both threads may read an old value of the total_balance at the same time due to context switching and then update it at almost the same time. As a result, the increment of the balance from one of the threads is lost. Also, on running it multiple times, we can see that value of total_balance keeps changing.

Part 2: Synchronization

To fix the issue of synchronization faced in last case, we have to use **spinlocks** that will allow update to happen atomically. To do so, we have to create a lock data structure called **thread_spinlock** and declare three functions:

1. **void thread_spin_init(struct thread_spinlock *lk)** to initialize the lock.
2. **void thread_spin_lock(struct thread_spinlock *lk)** to acquire the lock.
3. **void thread_spin_unlock(struct thread_spinlock *lk)** to release the lock.

```

3  struct thread_spinlock{
4      volatile uint lock;
5  };
6
7  void thread_spin_init(struct thread_spinlock *lk) {
8      lk->lock = 0;
9  };
10
11 void thread_spin_lock(struct thread_spinlock *lk) {
12     while (xchg(&lk->lock, 1) != 0)
13         ;
14     __sync_synchronize();
15 }
16
17 void thread_spin_unlock(struct thread_spinlock *lk) {
18     __sync_synchronize();
19     asm volatile("movl $0, %0" : "+m" (lk->lock) : );
20 }
21

```

- The **thread_spin_init()** function just initializes the value of variable **lock** to 0.

To acquire the lock, we used **xchg** function present in **x86.h** file which is basically a **compare and swap** assembly level instruction. It returns the current value of lock and if it is 0, it updates it to 1 but returns old value only.

The **Volatile** keyword is intended to prevent the compiler from applying any optimizations on objects that can change in ways that cannot be determined by the compiler.

```

120 static inline uint
121 xchg(volatile uint *addr, uint newval)
122 {
123     uint result;
124
125     // The + in "+m" denotes a read-modify-write operand.
126     asm volatile("lock; xchgl %0, %1" :
127         "+m" (*addr), "=a" (result) :
128         "1" (newval) :
129         "cc");
130     return result;
131 }

```

__sync_synchronise() function is used to synchronize data in all threads, i.e., the value of lock is updated in all threads simultaneously.

- The **thread_spin_lock()** function will check if lock is acquired or not. If not, it will acquire the lock and synchronize it across all threads.

If lock == 1, it means it is acquired by some thread, else if lock == 0, it is ready to be acquired by some thread.

- The **thread_spin_unlock()** function releases the lock by making its value equal to 0 via assembly level code.

After this, we copy *thread.c* to use spinlock around the critical section of code to new file called *thread_spinlock.c*. And similarly to last part, update *Makefile* to add this in command line instruction of Xv6 kernel.

```
thread_spin_lock(&lock);
old = total_balance;
delay(100000);
total_balance = old + 1;
thread_spin_unlock(&lock);
```

First, acquire the lock, then read the old value, increment it and write back the update value, thereafter releasing the lock.

```
$ thread_spinlock
SStatarrttiinnngg do_dwoork_: wso:rkb:2 s
:b1
Done s:2F8C
Done s:2F68
Threads finished: (4):5, (5):4, shared balance:6000
$ thread_spinlock
SStatarrtiinggt dino_g wdoor_kw:o sr:k:b 1s
:b2
Done s:2F8C
Done s:2F68
Threads finished: (7):8, (8):7, shared balance:6000
$ thread_spinlock
SStatarrttiinnngg d o_dwoor_wko:rk s:: bs2
:b1
Done s:2F8C
Done s:2F68
Threads finished: (10):11, (11):10, shared balance:6000
```

Upon execution of *thread_spinlock* in Xv6, we can see that no matter how many times we run it, the **value of total_balance is 6000** that is equal to expected one.

This happens because spinlock does not allow another thread to enter critical section and read/update the value of *total_balance* unless the lock is released. Only after releasing the lock, the second thread will be able to read/update this.

Hence, increment from both threads are reflected in final value.

Spinlocks that we have implemented might be inefficient in some cases. When all threads of the process run in parallel on different CPUs, spinlocks are perfect. However, if system has single physical CPU or it is under high load and a context switch occurs in a critical section, then all threads of the process start to spin endlessly, waiting for the lock-holding thread to be release lock.

One possible approach is to implement a different synchronization primitive, a **mutex**.

Mutex is similar to spinlock except that **spinlock causes a thread trying to acquire it to simply wait in the loop and repeatedly check for its availability** while **mutex is a program object that is created so that multiple processes can take turns sharing the same resource**, i.e., the CPU time is not wasted checking for the condition of while loop as process goes to sleep.

To do so, we have to create a lock data structure called *thread_mutex* and declare three functions:

1. *void thread_mutex_init(struct thread_mutex *m)* to initialize the lock.
2. *void thread_mutex_lock(struct thread_mutex *m)* to acquire the lock.
3. *void thread_mutex_unlock(struct thread_mutex *m)* to release the lock.

```
23 struct thread_mutex{
24     volatile uint lock;
25 };|
26
27 void thread_mutex_init(struct thread_mutex *m) {
28     m->lock = 0;
29 };
30
31 void thread_mutex_lock(struct thread_mutex *m) {
32     while (xchg(&m->lock, 1) != 0)
33         sleep(1);
34     __sync_synchronize();
35 };
36
37 void thread_mutex_unlock(struct thread_mutex *m) {
38     __sync_synchronize();
39     asm volatile("movl $0, %0" : "+m" (m->lock) : );
40 };
```

The implementation of this is exactly similar to that of spinlock except for **thread_mutex_lock** function in which we add *sleep(1)* in the while loop (line 33).

This allows the process to go in sleep if lock is not available so that the CPU can meanwhile finish other tasks instead of continuously checking for while loop condition.

```

thread_mutex_lock(&ml);
old = total_balance;
delay(100000);
total_balance = old + 1;
thread_mutex_unlock(&ml);

```

After this, we created a new file **thread_mutex.c** that uses mutex instead of spinlock in the test code and updated **Makefile** to add this as another command line instruction.

```

$ thread_mutex
SSttaarrtitingn gd od_o_wworork:k: ss::bb21

Done s:2F8C
Done s:2F68
Threads finished: (5):6, (6):5, shared balance:6000
$ thread_mutex
SSttaratritnign g ddoo__wwoorrk:: s s::bb21

Done s:2F68
Done s:2F8C
Threads finished: (8):8, (9):9, shared balance:6000

```

As expected, this time also, we got expected results, i.e., the final value of **total_balance** is 6000.

Key Differences between Spinlock and Mutex:

Spinlock	Mutex
It is a type of lock that causes a thread attempting to obtain it to check for its availability while waiting in a loop continuously.	It is a program object designed to allow different processes to take turns sharing the same resource.
The process may not sleep while waiting for the lock.	The process may sleep while waiting for the lock.
Spinlock makes no use of context switching.	It involves context switching.
It temporarily prevents a thread from moving.	It may block a thread for an extended amount of time.
It is useful for limited critical sections; else, it wastes the CPU cycles.	It is useful for crucial extended areas where frequent context switching would add overhead.
It does not support preemption.	It supports preemption.

Submission:

We created patch file using the command: *diff -ruN xv6-public xv6-public-assign1 > patchfile.patch*
 The submitted zip folder **C21.zip** contains:

1. patchfile.patch
2. Report.pdf
3. xv6-public-assign1: The xv6 directory containing the complete code for the assignment in-case the patch file does not work.