

CS_344 Assignment 2

Group C21: Satvik Tiwari, 200101091

Pranjal Baranwal, 200101083

Akshat Mittal, 200101011

Part A: Implementing System Calls

In this part, we have to implement following system calls in xv6:

1. *getNumProc()*

Returns the total number of active processes in the system (either in embryo, running, runnable, sleeping, or zombie states).
To do so, we simply looped over the process table to get all the active processes i.e. any process which is not in *UNUSED* state and returned the count.

```
592 // get total number of active processes in system
593 int
594 getNumProc(void)
595 {
596     struct proc *p;
597
598     int count = 0;
599     acquire(&ptable.lock);
600     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
601         if(p->state == UNUSED)
602             continue;
603         else
604             count++;
605         // cprintf("%d\n", p->pid);
606     }
607     release(&ptable.lock);
608     return count;
609 }
```

2. *getMaxPid()*

Returns the maximum process ID (*pid*) out of all currently active processes in the system. We simply looped over the process table to get the maximum *pid*.

```
611 // get the maximum PID among the PIDs of all currently active processes
612 int
613 getMaxPid(void)
614 {
615     struct proc *p;
616
617     int max = -1;
618     acquire(&ptable.lock);
619     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
620         if(p->state == UNUSED)
621             continue;
622         if(p->pid > max)
623             max = p->pid;
624     }
625     release(&ptable.lock);
626     return max;
627 }
```

3. `getProcInfo(pid, &processInfo)`

This function takes two arguments, ID of the process and a *processInfo struct*. It gives the *parentID*, the number of times context switched and the process size in bytes of the process with given ID.

For this, we added *contextswitches* field in our process structure in *proc.h* and simply looped over the process table. Upon finding the given ID, copy all the required data into our *processInfo struct*.

It returns -1 if this ID is not found in the process table.

```
// get info about the process having given PID
int
getProcInfo(int pid, struct processInfo* st)
{
    struct proc *p;
    int flag = -1;
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->pid == pid)
        {
            st->ppid = 0;
            // check if parent exists
            if(p->parent != 0)
                st->ppid = p->parent->pid;
            st->psize = p->sz;
            st->numberContextSwitches = p->contextswitches;
            flag = 0;
            break;
        }
    }
    release(&ptable.lock);
    return flag;
}
```

4. `set_burst_time(n)`

This function takes an integer input n and set the burst time of current process equal to n.

For this, we added burst field in our process structure. Using *myproc()* function we simply get instance to our current process and then set the burst time. We used *yield* to skip one CPU scheduling as scheduling might change due to this change.

```
653 // set burst time for a process
654 int
655 set_burst_time(int bt)
656 {
657     myproc()->burst = bt;
658     // cprintf("%d\n", myproc()->pid);
659     // skip one CPU scheduling round.
660     yield();
661     return 0;
662 }
```

5. `get_burst_time()`

It simply returns the burst time of current process.

```
664 int
665 get_burst_time(void)
666 {
667     return myproc()->burst;
668 }
```

Files updated to add system calls and user programs:

1. *proc.c* - All the above codes are written in this file.
2. *proc.h* - The proc structure is modified to include fields *contextswitches* and *burst*.

```
int contextswitches; // store number of context switches
int burst;           // store approximate burst time
```

3. *defs.h* - This header file is included in *sysproc.c* to allow it to call the functions implemented in *proc.c*, so we declare our functions here.

```
// assignmnet 2A
int      getNumProc(void);
int      getMaxPid(void);
int      getProcInfo(int, struct processInfo*);
int      set_burst_time(int);
int      get_burst_time(void);
```

4. **sysproc.c** - We need to pick the arguments provided by the user from stack using argptr, since the process related to system call must have a void argument. Here we call the actual functions defined in **proc.c**.

```
// assignment 2A
int
sys_getNumProc(void)
{
    return getNumProc();
}

int
sys_getMaxPid(void)
{
    return getMaxPid();
}

int
sys_getProcInfo(void)
{
    int pid;
    struct processInfo* st;
    if(argint(0, &pid) < 0)
        return -1;
    if(argptr(1, (void*)&st, sizeof(st)) < 0)
        return -1;
    return getProcInfo(pid, st);
}
```

```
int
sys_set_burst_time()
{
    int bt;
    if(argint(0, &bt) < 0)
        return -1;
    if(bt <= 0)
        return -1;
    set_burst_time(bt);
    return 0;
}

int
sys_get_burst_time(void)
{
    return get_burst_time();
}
```

5. **syscall.h** - Map the system call names to system call numbers.

```
#define SYS_getNumProc 22
#define SYS_getMaxPid 23
#define SYS_getProcInfo 24
#define SYS_set_burst_time 25
#define SYS_get_burst_time 26
```

6. **syscall.c** - Add function pointers to the actual system call implementations and export it. The variables declared in **syscall.h** are used to index into the array of function pointers.

```
// assignment 2A
extern int sys_getNumProc(void);
extern int sys_getMaxPid(void);
extern int sys_getProcInfo(void);
extern int sys_set_burst_time(void);
extern int sys_get_burst_time(void);
```

```
// assignment 2A
[SYS_getNumProc] sys_getNumProc,
[SYS_getMaxPid] sys_getMaxPid,
[SYS_getProcInfo] sys_getProcInfo,
[SYS_set_burst_time] sys_set_burst_time,
[SYS_get_burst_time] sys_get_burst_time
```

7. **user.h** - System call definitions are added here to make them available to the user program.

```
// Assignment 2A
int getNumProc(void);
int getMaxPid(void);
int getProcInfo(int pid, struct processInfo*);
int set_burst_time(int n);
int get_burst_time(void);
```

8. **usys.S** - The list of system calls which we want to export by kernel are added here.

```
SYSCALL(getNumProc)
SYSCALL(getMaxPid)
SYSCALL(getProcInfo)
SYSCALL(set_burst_time)
SYSCALL(get_burst_time)
```

- 9. **Makefile** - Makefile needs to be edited before our user program is available for xv6 source code for compilation. We included `_<userprogram_name>\` in the UPROGS list to add this as a command line instruction in XV6 kernel which will execute corresponding test files.

- 10. We also added **`getNumProc.c`**, **`getMaxPid.c`**, **`getProcInfo.c`** and **`set_burst_time.c`** to test the system calls we made.

```
UPROGS=\
_cat\
_echo\
_forktest\
_grep\
_init\
_kill\
_ln\
_ls\
_mkdir\
_rm\
_sh\
_stressfs\
_usertests\
_wc\
_zombie\
_getNumProc\
_getMaxPid\
_getProcInfo\
_set_burst_time\
```

Part B: Scheduling

Shortest Job First (SJF) Scheduling

The SJF scheduling algorithm schedules the process with lowest burst time first. It is implemented in the `scheduler()` function in **`proc.c`** file.

- We initialize `burst` and `contextswitches` field in `allocproc()` function in **`proc.c`**.

```
p->contextswitches = 0;
p->burst = 0;
```

- As we wanted to simulate this on a single CPU, the `NCPU` parameter is changed to 1 in **`param.h`**.

```
#define NCPU 1 // maximum number of CPUs
```

- In **`trap.c`**, lines 107 to 109 are commented to turn off the preemption in scheduler as SJF is non-preemptive.

```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105
106 // // Comment it out for shortest job first scheduling algorithm
107 // if(myproc() && myproc()->state == RUNNING &&
108 //     tf->trapno == T_IRQ0+IRQ_TIMER)
109 //     yield();
```

- Changes in scheduler function of **`proc.c`** for SJF scheduling:

1. We iterate over all the processes which are `RUNNABLE` and chose the one with the smallest burst time for scheduling.
2. If there is no process then we simply continue. Otherwise, we context-switch to this chosen process.
3. On line 403 we increased number of context switches by 1 for current process.

- If we have n processes then our scheduler has a time complexity of $O(n)$.

```

361 // Shortest job first scheduling algorithm
362 while(1){
363     // Enable interrupts on this processor.
364     sti();
365     // To store the job with least burst time
366     struct proc *shortest_job = 0;
367
368     // Acquire process table lock
369     acquire(&ptable.lock);
370     // Find the job with least burst time
371     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
372     {
373         if (p->state == RUNNABLE)
374         {
375             if (!shortest_job)
376             {
377                 shortest_job = p;
378             }
379             else if (p->burst < shortest_job->burst)
380             {
381                 shortest_job = p;
382             }
383         }
384     }
385
386     if (!shortest_job)
387     {
388         release(&ptable.lock);
389         continue;
390     }

```

```

392 p = shortest_job;
393
394 // Switch to chosen process. It is the process's job
395 // to release ptable.lock and then reacquire it
396 // before jumping back to us.
397 c->proc = p;
398 switchuvm(p);
399 p->state = RUNNING;
400
401 swtch(&(c->scheduler), p->context);
402 // increment number of context switches
403 p->contextswitches = p->contextswitches + 1;
404 switchkvm();
405
406 // Process is done running for now.
407 // It should have changed its p->state before coming back.
408 c->proc = 0;
409 release(&ptable.lock);
410 }

```

Testcases SJF Scheduler:

- We have our testcase in **testSched1.c** and **testSched2.c** in which we fork 10 processes one by one with different burst times. We make 6 CPU bound processes and 4 I/O bound processes.

```

$ testSched1
(10) CPU Bound(931036809)      Burst Time: 10  Context Switches: 5
(9) CPU Bound(874574840)      Burst Time: 20  Context Switches: 10
(8) CPU Bound(745846863)      Burst Time: 30  Context Switches: 15
(7) CPU Bound(467847273)      Burst Time: 40  Context Switches: 20
(6) CPU Bound(1284847454)     Burst Time: 50  Context Switches: 25
(5) CPU Bound(1130519139)     Burst Time: 60  Context Switches: 34
(4) IO Bound   Burst Time: 70  Context Switches: 701
(3) IO Bound   Burst Time: 80  Context Switches: 801
(2) IO Bound   Burst Time: 90  Context Switches: 901
(1) IO Bound   Burst Time: 100 Context Switches: 1001
$
$
$
$
$ testSched2
(3) CPU Bound(491195140)      Burst Time: 10  Context Switches: 6
(6) CPU Bound(1872725030)     Burst Time: 30  Context Switches: 16
(5) CPU Bound(-2147483648)    Burst Time: 60  Context Switches: 34
(2) CPU Bound(487655483)      Burst Time: 70  Context Switches: 39
(8) CPU Bound(981612159)      Burst Time: 80  Context Switches: 45
(9) CPU Bound(-2147483648)    Burst Time: 100 Context Switches: 56
(7) IO Bound   Burst Time: 20  Context Switches: 201
(1) IO Bound   Burst Time: 40  Context Switches: 401
(10) IO Bound  Burst Time: 50  Context Switches: 501
(4) IO Bound   Burst Time: 90  Context Switches: 901

```

Fig- Output of Round Robin scheduling

```

$ testSched1
(10) CPU Bound(174984499)      Burst Time: 10  Context Switches: 1
(9) CPU Bound(349968999)      Burst Time: 20  Context Switches: 1
(8) CPU Bound(524953499)      Burst Time: 30  Context Switches: 1
(7) CPU Bound(699938000)      Burst Time: 40  Context Switches: 1
(6) CPU Bound(874922500)      Burst Time: 50  Context Switches: 1
(5) CPU Bound(1049907000)     Burst Time: 60  Context Switches: 1
(4) IO Bound   Burst Time: 70  Context Switches: 701
(3) IO Bound   Burst Time: 80  Context Switches: 801
(2) IO Bound   Burst Time: 90  Context Switches: 901
(1) IO Bound   Burst Time: 100 Context Switches: 1001
$
$
$
$
$ testSched2
(3) CPU Bound(174984499)      Burst Time: 10  Context Switches: 1
(6) CPU Bound(524953499)      Burst Time: 30  Context Switches: 1
(5) CPU Bound(1049907000)     Burst Time: 60  Context Switches: 1
(2) CPU Bound(1224891500)     Burst Time: 70  Context Switches: 1
(8) CPU Bound(1399876000)     Burst Time: 80  Context Switches: 1
(9) CPU Bound(1749845000)     Burst Time: 100 Context Switches: 1
(7) IO Bound   Burst Time: 20  Context Switches: 201
(1) IO Bound   Burst Time: 40  Context Switches: 401
(10) IO Bound  Burst Time: 50  Context Switches: 501
(4) IO Bound   Burst Time: 90  Context Switches: 901

```

Fig- Output of SJF scheduling

- For testcase1, we have forked 10 processes. The first 6 are CPU bound processes with burst time of 100, 90, 80, 70, 60 and 50 units respectively while the last 4 are IO bound process having burst time of 40, 30, 20 and 10 units respectively.
- For testcase2, we have forked 10 processes and for every IO bound process, we fork two CPU bound processes except the 4th IO bound process (10th process). The burst time of the processes are: 40 (IO), 70 (CPU), 10 (CPU), 90 (IO), 60 (CPU), 30 (CPU), 20 (IO), 80 (CPU), 90 (CPU) and 50 (IO).

- In normal round robin, we can see that the order of completion depends on burst time, shorter jobs are getting completed first but there are several preemptions for all processes before they complete. IO processes took more time to complete (though they have less burst time). This is due to a greater number of IO operations; they are pushed from running to waiting state and thus, the number of context switches also increases.
- While in SJF, we do not preempt the process until completed and hence only 1 context switch is observed in case of CPU bound processes while context switches still occur in IO bound processes due to IO operation.
- When we execute testcase1 and testcase2, we see that all CPU bound processes terminate before IO bound processes and both CPU bound and IO bound processes terminate in ascending order of burst times among themselves, showcasing a successful SJF scheduling algorithm. IO bound processes have a large number of context switches because of lots of IO delays.

Hybrid Scheduler (SJF + Round Robin)

The hybrid scheduling algorithm is similar to round-robin but the processes are selected not on basis of first-come-first serve but on shortest burst time. It is implemented in the `scheduler()` function in ***proc.c*** file.

- The ***struct proc*** was also modified wherein additional members such as `time_slice` and `first_proc` were included to keep track of the time slice taken by a process and the shortest process so as to change the `time_quanta` variable as per the `time_slice` required for the `first_proc`, i.e., the shortest burst time process.

```
int time_slice; // for time quanta
int first_proc; // to indicate shortest process
```

- We initialize burst, contextswitches, `time_slice` and `first_proc` field in `allocproc()` function in ***proc.c***

```
98 // For hybrid scheduling
99 p->burst = 1;
100 p->time_slice = 1;
101 p->first_proc = 0;
```

- In ***trap.c***, for round-robin, the preemption happens when time-quantum expires for a process. But for hybrid scheduling, we have to consider the burst time as well since we are using SJF.

Hence, for the `first_proc`, i.e., the shortest process, we increment the `time slice` and update `time_quanta` for the round.

For all remaining processes, if they have not run for `time_quanta` amount of time, we simply increment the `time_slice`, else as soon as their `time_quanta` expire, we preempt the current process.

```
109 if(myproc() && myproc()->state == RUNNING &&
110    tf->trapno == T_IRQ0+IRQ_TIMER)
111 {
112     // yield();
113     // For hybrid scheduling algorithm
114     if(myproc()->first_proc && (first_pid == -1 || first_pid == myproc()->pid))
115     {
116         myproc()->time_slice++;
117         time_quanta = myproc()->time_slice + 1;
118         first_pid = myproc()->pid;
119     }
120     else
121     {
122         if(myproc()->time_slice < time_quanta)
123         {
124             myproc()->time_slice++;
125         }
126         else {
127             myproc()->time_slice = 0;
128             yield();
129         }
130     }
131 }
```

- Changes in scheduler function of ***proc.c*** for hybrid scheduling:

1. First, a ready queue is created wherein all the `RUNNABLE` processes will be pushed.
2. Then, we sort the processes in ready queue using standard **merge sort algorithm** based on their burst time and mark first process of queue as `first_proc`.

- Then from line 397 to 421, we schedule the processes one by one with increasing burst time and doing round robin at the same time.

- If we have n processes then our scheduler has a time complexity of $O(n \cdot \log n)$ which is time complexity to sort the processes.

```

369 // Hybrid scheduling algorithm
370 int flag = 1;
371 while(1){
372     // Enable interrupts on this processor.
373     sti();
374
375     // Acquire process table lock
376     acquire(&ptable.lock);
377     // Set up Ready Queue
378     struct proc* ready_queue[NPROC];
379     int k = 0;
380     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
381     {
382         if(p->state == RUNNABLE)
383         {
384             ready_queue[k++] = p;
385         }
386     }
387
388     // Sort Ready Queue
389     merge_sort(ready_queue, 0, k-1);
390
391     if(k && flag)
392     {
393         ready_queue[0]->first_proc = 1;
394         flag = 0;
395     }

```

```

396
397 // Schedule the jobs with increasing burst time
398 for (int i = 0; i <= k-1; i++)
399 {
400     p = ready_queue[i];
401     if (p->state == RUNNABLE)
402     {
403         // Switch to chosen process. It is the process's job
404         // to release ptable.lock and then reacquire it
405         // before jumping back to us.
406         c->proc = p;
407         switchvm(p);
408         p->state = RUNNING;
409
410         swtch(&(c->scheduler), p->context);
411
412         // increment number of context switches
413         p->contextswitches = p->contextswitches + 1;
414
415         switchkvm();
416
417         // Process is done running for now.
418         // It should have changed its p->state before coming back.
419         c->proc = 0;
420     }
421 }
422 release(&ptable.lock);
423

```

Testcases Hybrid Scheduler:

- We have our testcase in **testSched1.c** and **testSched2.c** in which we fork 10 processes one by one with different burst times. We make 6 CPU bound processes and 4 I/O bound processes.

```

$ testSched1
(10) CPU Bound(663776823)      Burst Time: 10  Context Switches: 2
(9) CPU Bound(966320054)      Burst Time: 20  Context Switches: 3
(8) CPU Bound(818218982)      Burst Time: 30  Context Switches: 4
(7) CPU Bound(241819969)      Burst Time: 40  Context Switches: 5
(6) CPU Bound(1242780935)     Burst Time: 50  Context Switches: 7
(5) CPU Bound(1166756679)     Burst Time: 60  Context Switches: 8
(4) IO Bound      Burst Time: 70  Context Switches: 701
(3) IO Bound      Burst Time: 80  Context Switches: 801
(2) IO Bound      Burst Time: 90  Context Switches: 901
(1) IO Bound      Burst Time: 100 Context Switches: 1001
$
$
$
$
$ testSched2
(3) CPU Bound(203214434)      Burst Time: 10  Context Switches: 2
(6) CPU Bound(1015743411)     Burst Time: 30  Context Switches: 4
(5) CPU Bound(-2147483648)     Burst Time: 60  Context Switches: 8
(2) CPU Bound(1068254894)     Burst Time: 70  Context Switches: 9
(8) CPU Bound(803224280)      Burst Time: 80  Context Switches: 10
(9) CPU Bound(1454928048)     Burst Time: 100 Context Switches: 13
(7) IO Bound      Burst Time: 20  Context Switches: 201
(1) IO Bound      Burst Time: 40  Context Switches: 401
(10) IO Bound     Burst Time: 50  Context Switches: 501
(4) IO Bound      Burst Time: 90  Context Switches: 901

```

Fig- Output of Round Robin scheduling

- For testcase1, we have forked 10 processes.
The first 6 are CPU bound processes with burst time of 100, 90, 80, 70, 60 and 50 units respectively while the last 4 are IO bound process having burst time of 40, 30, 20 and 10 units respectively.
 - For testcase2, we have forked 10 processes and for every IO bound process, we fork two CPU bound processes except the 4th IO bound process (10th process).
The burst time of the processes are: 40 (IO), 70 (CPU), 10 (CPU), 90 (IO), 60 (CPU), 30 (CPU), 20 (IO), 80 (CPU), 90 (CPU) and 50 (IO).
 - When we execute testcase1 and testcase2, we see that all CPU bound processes terminate before IO bound processes and both CPU bound processes and I/O bound processes terminate in ascending order of burst times among themselves, showcasing a successful hybrid scheduling algorithm.
The only difference between hybrid scheduler and SJF is that SJF was non preemptive but hybrid scheduler is preemptive and hence we can see some context switches for CPU bound process too unlike SJF.
IO bound processes have a large number of context switches because of lots of IO delays.
-

Submission:

C21.zip contains following:

1. **Report.pdf**
2. **partAB** – Folder containing all modified files for part A and SJF scheduling
3. **partAB.patch**
4. **hybrid** – Folder containing all modified files for hybrid scheduling
5. **hybrid.patch**
6. **xv6-public-assign2-sjfs** – xv6 folder corresponding to partAB.patch
7. **xv6-public-assign2-hybrid** – xv6 folder corresponding to hybrid.patch